

On the Role of Chain of Thoughts in the Long In-Context Learning

Anonymous ACL submission

001

Abstract

In-Context Learning (ICL) has emerged as a powerful paradigm for adapting Large Language Models (LLMs) to new tasks without gradient updates. While advances in long-context models have enabled a shift from few-shot to many-shot ICL, achieving performance comparable to fine-tuning, research has largely focused on classification tasks. Therefore we target to address the underexplored behavior of Chain-of-Thought (CoT) prompting in many-shot scenarios, explaining the shift from studies in how to select appropriate ICL examples to how to enable LLMs to evolve at test-time. We present a comprehensive analysis of many-shot in-context CoT learning, uncovering behavioral differences between reasoning-oriented and non-reasoning oriented LLMs. Our findings show that with both types of models, there is a fundamental difference from earlier studies in the ICL setting. Demonstration selection based on input similarity, which is a common heuristic in ICL, becomes ineffective under the CoT paradigm. Counterintuitively, our experiments show that the quality of the reasoning chain, as measured by its ground-truth correctness, is not the primary factor for success. Instead, we observe a consistent performance gap where model-self-generated CoTs with incorrectness outperform those with human-verified, correct reasoning. This suggests that the success of many-shot CoT prompting is driven more by alignment with the model's internal reasoning patterns than by the objective quality of the demonstrations. These observations indicate that LLMs are more effective at learning from their own "experiences" than from ground-truth reasoning or externally constructed examples, which they may find harder to interpret. This insight also explains why similarity-based selection strategies lose their effectiveness in CoT settings. Additionally, our experiments reveal that, unlike previous many-shot ICL studies, the order of CoT demonstrations significantly affects performance. We observe a fluctuating or increasing variance along the number of shots when the demonstration order is shuffled. This points to the critical role of ordering in the effectiveness of LLM test-time learning.

1 Introduction

In-Context Learning (ICL), where Large Language Models (LLMs) are prompted with a sequence of input-output demonstrations and asked to produce predictions for new inputs without any gradient updates, has gained significant attention. Research has extensively investigated its benefits (Sorensen et al., 2022; An et al., 2023; Mavromatis et al., 2023) and underlying mechanisms (Min et al., 2022; von Oswald et al., 2023; Deutch et al., 2024). A substantial body of work has focused on enhancing ICL, including further pre-training models for improved ICL capability, developing strategies for demonstration selection, and exploring the transition from few-shot to many-shot learning. While early work focused on few-shot ICL, recent advances in scaling context windows have made it possible to explore many-shot ICL, where dozens to hundreds or even more demonstrations can be provided, allowing performance comparable to fine-tuning (Agarwal et al., 2024; Bertsch et al., 2025; Baek et al., 2025). However, the majority of these studies focus on classification-oriented ICL, where tasks are relatively shallow and answers are directly inferred from patterns in the demonstrations.

In parallel, there has been growing interest in chain-of-thought (CoT) prompting, a technique that improves reasoning tasks by encouraging models to generate intermediate reasoning steps before arriving at a final answer (Kojima et al., 2022). While CoT has shown strong performance in few-shot settings (Zhang et al., 2023; Wei et al., 2022; Luo et al., 2023), its behavior in many-shot in-context learning remains underexplored. Crucially, in the case of ICL with CoT, most prior work directly applies the few-shot paradigm without delving deeply

086 into the underlying mechanisms. This raises several critical questions.
087

- 088 1. What happens in the many-shot scenario: does
089 performance scale monotonically, or does it
090 plateau or even degrade?
091 2. Is in-context CoT fundamentally different
092 from ICL with a single label?
093 3. What correlating factors govern the effective-
094 ness of many-shot CoT, and can previous
095 demonstration selection strategies (e.g., based
096 on semantic similarity) be applied directly?

097 This gap is significant in light of the growing
098 context lengths supported by LLMs and the emerging
099 concept of test-time scaling and DeepResearch.
100 Scaling in test-time enables LLMs to refine their
101 responses without parameter updates during in-
102 ference, through paradigms such as sequential re-
103 vision and parallel sampling (Snell et al., 2025).
104 Prior works have explored inference-time align-
105 ment methods with ICL (Lin et al., 2024; Li et al.,
106 2025). Li et al. (2025) proposes parallel sampling
107 for sequential refining and shows that increasing
108 the search width of sampled responses in in-context
109 consistently enhances the performance, demon-
110 strating the potential of incorporating in-context
111 demonstrations in enhancing the LLM capability
112 during inference. Yet it remains unclear whether
113 in-context CoT similarly benefits from longer con-
114 texts, or whether it introduces new challenges due
115 to the complexity of reasoning chains.

116 In this work, we present a comprehensive anal-
117 ysis of many-shot in-context CoT learning, com-
118 paring its behavior with traditional classifica-
119 tion-based ICL and evaluating its effectiveness under ex-
120 tended context lengths. Our study identifies several
121 fundamental differences between CoT and classi-
122 fication ICL under many-shot settings. We find
123 that, in contrast to many-shot ICL in tasks without
124 involving CoT (Bertsch et al., 2025; Baek et al.,
125 2025), the performance of in-context CoT is highly
126 sensitive to demonstration ordering and selection.
127 For demonstration selection, unlike traditional ICL,
128 similarity no longer serves as a reliable signal, high-
129 lighting the need for further investigation specifi-
130 cally for in-context CoT. These findings suggest
131 that many-shot CoT learning is governed by differ-
132 ent dynamics than those observed in traditional ICL.
133 To address these challenges, we investigate factors
134 correlated with in-context CoT performance and

135 find that constructing LLM-aligned CoT demon-
136 strations stabilizes its performance. Our results
137 show that the quality of the provided CoT is unex-
138 pectedly not a critical factor. We observe a clear hi-
139 erarchy where demonstrations with self-generated,
140 incorrect CoTs lead to the best performance outper-
141 forming ground-truth human-verified CoTs. Adopt-
142 ing CoTs generated by a stronger, more advanced
143 LLM, on the other hand, results in a lower accuracy.
144 This suggests that the mechanism behind many-
145 shot CoT is not merely about providing higher-
146 quality examples, but is related to the alignment to
147 LLM.

2 Mysteries of many-shot ICL

2.1 Experiment Setup

Tasks. Previous studies in many-shot (Li et al.,
150 2024; Bertsch et al., 2025) lacks exploration
151 in the reasoning tasks. Experiments are con-
152 ducted with a diverse type of tasks, includ-
153 ing traditional classification tasks (i.e., Super-
154 GLUE (Wang et al., 2019) with a narrow label
155 space; NLU (nlu, 2021), TREC (Hovy et al., 2001)
156 and BANKING77 (Casanueva et al., 2020) with
157 significantly larger label space), mathematical rea-
158 soning (i.e., GSM8K (Cobbe et al., 2021) and
159 MATH (Hendrycks et al., 2021)).

ICL Settings. For the reasoning tasks, the in-
161 clusion of CoT is natural. Hoping to provide the
162 complete picture also for the classification tasks,
163 we also adopt the LLM-generated CoT for the clas-
164 sification tasks.

- **Traditional ICL.** An instance consists of an input-output pair (x, y) . With n in-context demon-
166 strations provided, the LLM processes
167 the input as $\text{LLM}(x' \mid \{(x_i, y_i)\}_{i=1}^n)$ to gen-
168 erate y' . Since the token length of individ-
169 ual instances is relatively small, current open-
170 source LLMs can easily handle hundreds or
171 even thousands of examples during evalua-
172 tion.

- **In-context CoT.** An instance consists of an input-CoT-output triplet (x, C, y) . With n in-context demon-
175 strations provided, the LLM processes the input as $\text{LLM}(x' \mid \{(x_i, C_i, y_i)\}_{i=1}^n)$ to generate y' . Since the
176 token length of individual instances can be
177 significantly larger, depending on the token
178 length of C , current open-source LLMs are

183 typically limited to evaluating only around
 184 hundreds of examples.

185 **LLMs Studied.** Building on previous studies that
 186 highlight the limitations of long-context LLMs
 187 in ICL (Li et al., 2024; Bertsch et al., 2025)
 188 and given the recent advancements in instruction-
 189 tuned models with extended context windows in
 190 about 130K tokens, we conduct experiments with
 191 LLaMA 3.1 (Llama-3.1-8B-Instruct), LLaMA 3.3
 192 (Llama-3.3-70B-Instruct) (MetaAI, 2024), Qwen
 193 2.5 (7B) (Qwen2.5-7B-Instruct), Qwen 2.5 (14B)
 194 (Qwen2.5-14B-Instruct), Qwen 3 (8B) (Qwen3-8B),
 195 Qwen 3 (14B) (Qwen3-14B) (Qwen et al., 2025),
 196 enabling analysis across different LLM architec-
 197 tures and model sizes. To enable the processing of
 198 long context to the 131k token level for the Qwen
 199 family, we modified the config file and add the
 200 `rope_scaling` fields.


201 Unlike prior studies that focus primarily on clas-
 202 sification tasks with constrained decoding (Bertsch
 203 et al., 2025), we adopt a generative framework for
 204 both classification and generation tasks. Specif-
 205 ically, we formulate all tasks as text generation
 206 problems and evaluate model outputs using exact
 207 match against reference answers or labels. This
 208 approach aligns with recent trends in LLM research,
 209 where the emphasis has increasingly shifted toward
 210 open-ended generation settings.

211 In many-shot in-context CoT learning, the num-
 212 ber of tokens per demonstration can be substan-
 213 tially larger than in traditional many-shot ICL due
 214 to the length of CoT reasoning. For instance, when
 215 comparing geometry task to BANKING77, the aver-
 216 age demonstration length in the former is 30
 217 times longer, averaged across the training set. To
 218 maintain consistency with in-context CoT scaling,
 219 we prompt with approximately 100-shot in-context
 220 demonstrations in the following studies.


221 In addition, empirical results in the following
 222 subsection show that even with the Qwen3 family,
 223 model performance declines sharply beyond a cer-
 224 tain number of tokens. This suggests limited ben-
 225 efit in further increasing the number of in-context
 226 demonstrations under current model constraints.
 227 Under these considerations, our analysis in per-
 228 formed on the scope of about a hundred demon-
 229 strations.

230 2.2 Results

231 To enable a direct comparison of performance
 232 across tasks with varying accuracy ranges, we nor-




233 Figure 1: Performance comparison using normalized
 234 results between classification tasks (in **warm colors**) and
 235 math reasoning tasks (in **cool colors**). The x-axis repre-
 236 senters normalized accuracy, while the y-axis indicates
 237 the number of in-context demonstrations.




238 Figure 2: Performance on randomly sampled subset
 239 of math reasoning tasks using Llama-3.3-70B-Instruct
 240 (left) and QwQ-32B (right).

233 malize all results, as illustrated in Figure 1. De-
 234 tailed results are provided in Appendix B.

235 As shown in the figure, there is a significant
 236 difference between the two types of tasks: classifi-
 237 cation tasks exhibit a consistent pattern of steady
 238 improvement as the number of demonstrations in-
 239 creases, whereas math reasoning tasks show fluc-
 240 tuating or even declining performance. As pointed



(a) Qwen3-8B



(b) Qwen3-14B

Figure 3: Performance on randomly sampled subset of math reasoning tasks using Qwen 3 family.

out by Li et al. (2024), LLMs often struggle to learn effectively from long in-context examples.

With the advancement of LLMs, subsequent studies have shown that learning from long-context classification tasks has become increasingly effective (Agarwal et al., 2024; Bertsch et al., 2025). While LLMs demonstrate strong many-shot learning capabilities in classification tasks, their performance in reasoning-intensive tasks using in-context chain-of-thought (CoT) remains limited. This limitation persists even in recent LLMs: both the Qwen 2.5 and LLaMA 3.1 families show difficulty in handling long-context CoT settings. To further investigate, we evaluate a larger model, LLaMA-3.3-70B. As shown in Figure 2, even with this increased parameter size, the model is still unable to effectively learn from the provided demonstrations.

The recently released Qwen 3 family demonstrates a general trend of steadily increasing performance, as shown in Figure 3. To explore the limits of their capability, we leveraged its extended context window, pushing to the maximum token limit of 131k. This allowed for the inclusion of up to 256 in-context demonstrations in the geometry task, with each demonstration comprising approximately 450 tokens. However, we observe a severe and progressive performance degradation beyond a certain threshold of demonstrations. For instance, on the geometry task, the accuracy of Qwen3-14B progressively degrades from 64.92% (with 128 demonstrations) to 55.74% (with 181




Figure 4: The original performance with the error band computed across five random orders shuffled with a unified set of random seeds under Qwen 2.5 (14B). geometry* is evaluated with another randomly sampled in-context CoT demonstration set.

demonstrations), before collapsing to 12.32% (with 256 demonstrations). A similar pattern is observed with LLaMA 3.1 and Qwen 2.5 (8B), while Qwen 2.5 (14B) exhibits a less severe, though still present, fluctuation in performance.

Consequently, despite the Qwen 3 family’s robust overall performance, our experiments show that most LLMs struggle to effectively learn from a large number of in-context Chain-of-Thought demonstrations. This consistent failure mode across model families raises a research question: can any strategies be adopted to enhance CoT performance in long-context reasoning tasks?

2.3 Is LLM truly incapable to be benefitted from in-context CoT demonstrations?

While the recently released models no longer struggle with the large space classification tasks assessed by Li et al. (2024), new challenges arrived with the inclusion of the reasoning chain. To further investigate the effectiveness of in-context CoT, we conduct a focused analysis on three tasks (i.e., GSM8K, geometry and number_theory) with Qwen 2.5 (14B). Our analysis proceeds along two

295 strategies:

- 296 1. **Demonstration Order Sensitivity:** We ran-
297 domly shuffle the order of demonstrations five
298 times to examine whether an ordering that fa-
299 cilitates better in-context CoT exists.

- 300 2. **Demonstration Choice Sensitivity:** We ran-
301 domly sample alternative sets of in-context
302 examples to assess whether performance im-
303 provements can be attributed to the choice of
304 demonstrations.

305 The results are shown in Figure 4, presenting
306 the original performance trends and the error band
307 of mean±standard deviation across five accuracies
308 with shuffled orders. To ensure experimental fair-
309 ness, a fixed set of seeds is randomly sampled to
310 perform the order shuffles. The same seeds are con-
311 sistently applied across all configurations, spanning
312 different LLMs, numbers of in-context examples,
313 and tasks to maintain controlled variability.

314 For the GSM8K and number_theory tasks, an
315 increasing trend is observed with the first strat-
316 egy. Notably, although the original trend for num-
317 ber_theory was decreasing, both the upper-bound
318 performance and the average performance across
319 five orderings show an increasing trend. These in-
320 dicate a sensitivity to the order of demonstrations.

321 In the geometry task, the same trend is observed
322 after the initial attempt of applying the second strat-
323 egy, continuing up to 64 demonstrations. The sub-
324 sequent decline in performance can be attributed to
325 the context length exceeding 40k tokens, which ne-
326 cessitates the use of RoPE scaling. This indicates a
327 sensitivity to the selection of demonstrations.

328 3 Do LLMs reason with in-context CoT?

329 A key question in analyzing the performance of
330 in-context CoT reasoning will be whether its suc-
331 cess directly correlates with the model’s task un-
332 derstanding or whether other factors play a more
333 significant role. To investigate this, we create
334 two 4-option multiple-choice question answering
335 (MCQA) tasks and found that in-context CoT may
336 correlate more with data distribution than task un-
337 derstanding in Appendix A.

338 3.1 Correlating Factor Influencing In-Context 339 CoT Success

340 **Tasks.** We follow the previous settings and se-
341 lect a subset of math reasoning tasks for further
342 analysis, including geometry, number theory, and

343 GSM8K. With wider application of LLMs into real-
344 life tasks, especially with the development of Deep
345 Research (Huang et al., 2025; Yu et al., 2025), the
346 rationale behind comprehensive reasoning tasks
347 also matters. Therefore, apart from the math prob-
348 lems and classification tasks, we additionally stud-
349 ied LLM performance that required narrative rea-
350 soning with DetectiveQA (Xu et al., 2025). Since
351 DetectiveQA provides the corresponding evidence
352 and the reasoning chain for deriving the answer
353 based on the evidence. For each instance, the ev-
354 idence is provided as part of the question. The
355 corresponding CoT will be the derivation labelled
356 with “-1”.

357 **LLMs and Tasks Studied** We evaluate with
358 LLaMA 3.1 and Qwen 2.5 (14B) on three math rea-
359 soning benchmarks (i.e., GSM8K, number_theory,
360 and geometry) and evaluate with Qwen 3 (8B)
361 and Qwen 3 (14B) on number_theory and Detec-
362 tiveQA.

363 **Model-CoT alignment.** We investigate whether
364 the efficacy of in-context Chain-of-Thought (CoT)
365 learning is more significantly influenced by the log-
366 ical quality of the reasoning or its alignment with
367 the LLM’s own generative distribution. To this
368 end, we generate CoT demonstrations by prompt-
369 ing the LLM on the training set, rather than using
370 the dataset’s ground-truth CoT. These generated
371 CoT are then used as the in-context examples dur-
372 ing evaluation over the test set.

373 To investigate whether CoT quality or distribu-
374 tional alignment plays a more significant role in
375 performance, we construct three distinct demon-
376 stration sets to isolate these factors:

- 377 1. The Correct Set (cr): Samples where the
378 model’s generated answer is correct.

- 379 2. The Incorrect Set (wr): Samples where the
380 model’s generated answer is incorrect.

- 381 3. The First Set (first): The initial generation for
382 each instance, regardless of accuracy.

383 Each LLM is prompted 10 times per training
384 instance using a temperature of 1.0 for diversity.
385 During the 10 times of prompting, if the predicted
386 answer is correct, we include the corresponding
387 CoT and generated answer in the cr set; otherwise,
388 it is placed in the wr set. The first generation is
389 included in the first set. These resulting demon-
390 stration sets are compared against the original CoT
391 (i.e., origin) provided within the datasets.




Figure 5: Performance of two sets of self-generated in-context CoT, including the set filtered with only correct answer(cr) and the set filtered with only wrong answer(wr). $cr_{qwen14b}$ is prompting the LLaMA model with the in-context CoT generated by Qwen 2.5 (14B).




Figure 6: Performance of the first set of self-generated in-context CoT. $cr_{qwen3(14b)}$ is prompting the Qwen 3 (8B) model with the in-context CoT generated by Qwen 3 (14B).

392
393
394
395
396
397
398
399
400
401
402
403
404
405
Due to the high accuracy of both LLMs on
GSM8K, it is difficult to obtain incorrect outputs
even at a high temperature. Thus, the wr set is only
constructed for number_theory and geometry.

Result Surprisingly, as illustrated in Figure 5, the wr set with wrong answers and presumably flawed reasoning always outperforms the original CoT and performs comparably to the cr set across both LLMs and both tasks. This shows the effectiveness of having LLM-aligned in-context CoT. Additionally, with the self-generated CoT, both LLMs suffer significantly less from the sudden drop and great fluctuation issues, especially for LLaMA 3.1.

Moreover, the use of self-generated CoT signif-

406
407
408
409
410
411
412
413
414
415
icantly mitigates the issues of performance instability and sudden accuracy drops observed with the origin, an effect particularly pronounced for LLaMA 3.1. Collectively, these results suggest that the distributional characteristics of the in-context examples (i.e., their alignment with the LLM’s own generative patterns) exert a more substantial influence on stable CoT prompting than the conventional metric of quality, defined here as the presence of a correct final answer.

416
417
418
419
420
In the meantime, since the wr and cr sets were
constructed by sampling model outputs, an instance
can have all correct or all incorrect answers across
all runs, risking an empty instance. To ensure ro-
bustness, we perform the analysis using the first set

421 with guaranteed no emptiness.


422 Results in figure 6 reinforce the initial finding.
423 The first set again outperforms the original
424 CoT in both the number theory and DetectiveQA
425 tasks, suggesting data distribution is a more influ-
426 ential factor than quality for stable in-context CoT
427 prompting.

428 **Does “better” CoT give better and stable per-
429 formance?** To further assess the role of CoT quality,
430 we investigate whether better CoT from a better-
431 performing model can improve the performance
432 of a weaker one. Specifically, we prompt LLaMA
433 3.1 using CoT generated by Qwen 2.5 (14B) in
434 Figure 5 and Qwen 3 (8B) using CoT generated
435 by Qwen 3 (14B) in Figure 6, where Qwen 2.5
436 (14B) and Qwen 3 (14B) both show a better or
437 comparable performance with the baseline.

438 As shown in the olive line in Figure 5 and 6,
439 while LLaMA 3.1 and Qwen 3 (8B) does bene-
440 fit from higher-quality CoT with a performance
441 increase in geometry and DetectiveQA at the be-
442 ginning and at certain shots of in-context examples,
443 the model still suffers significantly from instability,
444 including occasional sudden performance drops
445 and great fluctuations. This implies that while
446 higher-quality CoT may enhance performance, but
447 likely to result in greater instability without good
448 in-context data distribution (i.e., not well-aligned
449 with the evaluation LLM). This further reinforces
450 our finding that data distribution is a more crucial
451 factor in successful in-context CoT prompting.

452 3.2 Rethinking the role of similarity

453 Previous research in in-context learning has shown
454 that retrieving semantically similar examples often
455 enhances model performance (Liu et al., 2022; Wu
456 et al., 2023; Kapuriya et al., 2025). To further investi-
457 giate this claim, we include the BANKING77 clas-
458 sification task as a control experiment. Specifically,
459 we construct two unified sets of in-context exam-
460 ples: one comprising the most semantically similar
461 examples and the other comprising the most dissim-
462 ilar examples to the test set. Similarity is measured
463 by computing the cosine similarity between ques-
464 tion embeddings, averaged across the entire test
465 set with a sentence transformer (all-mnlp-base-v2).
466 Approximately 250 samples are retrieved from the
467 training set to form the candidate pool for con-
468 structing the similar and dissimilar example sets.
469 We evaluate performance on both the BANKING77
470 task and the reasoning tasks introduced in Section



471 Figure 7: Performance with similarity(sim) and dissim-
472 ilar(dis) sets with Qwen 2.5 (14B). The area between
473 the two sets is filled with colors, indicating the relative
474 performance at each point.

475 2.3 with Qwen 2.5 (14B).

476 The results are presented in Figure 7. Since our
477 retrieval strategy is based on the global similarity
478 to the full test set rather than per-instance similarity,
479 the benefits of similarity-based retrieval may not be
480 as apparent in few-shot ICL. However, in settings
481 with more than 20 in-context examples, the similar
482 set consistently yields a better performance over
483 the dissimilar in BANKING77, with the area in be-
484 tween highlighted in green. This aligns with prior
485 findings in classification-based ICL. In contrast, we
486 observe the opposite trend for the three reasoning
487 tasks. With the increasing number of in-context
488 CoT provided, the dissimilar set consistently out-
489 performs the similar set, showing prior findings in
490 ICL cannot be extended to in-context CoT.

491 4 Direction Towards Test-time Scaling.

492 With demonstrated effectiveness with self-
493 generated CoT, learning is demonstrated within
494 the provided instance. This creates the question of
495 how to enable better learning from the provided
496 demonstrations. Bertsch et al. (2025); Baek et al.
497 (2025) found that the impact of demonstration
498 ordering diminishes as the number of demon-
499 strations increases. Building on this insight, we
500 further explore order sensitivity in in-context
501 chain-of-thought (CoT) prompting. Specifically,




Figure 8: Standard deviation across five order sampling on Qwen 2.5 (14B) and Qwen 3 (14B).

we calculate the standard deviation across five accuracy scores obtained from randomly shuffled demonstration orders. Consistent with prior findings, classification tasks (e.g., NLU and BANKING77) show a clear pattern, the standard deviation in performance decreases as more demonstrations are added. Aligning to previous understanding, classification tasks can be resolved by locating similar instances, leading to the stability of performance after providing a sufficiently large ICL examples, no matter how the order of demonstrations is.

In contrast, for reasoning tasks, the standard deviation either fluctuates unpredictably or gradually increases with more demonstrations. This implies that, with complex problems necessitating reasoning, order is one of the important factors affecting how LLM can learn during test-time. This also align to human learning, where the chapters of the textbook are arranged in a specifically designed order to foster better learning and understanding.

5 Related Works

5.1 Many-shot ICL

With studies enabling LLMs to handle longer context lengths (Peng et al., 2024; Han et al., 2024; Ding et al., 2024), Agarwal et al. (2024) introduce the concept of many-shot ICL, which incorporates a significantly larger number of in-context demonstrations. Their results show a comparable performance to fine-tuning across various types of tasks. Subsequent studies (Li et al., 2024; Bertsch et al., 2025) investigate the effectiveness of many-shot ICL in open-source LLMs and examine its distinct characteristics compared to few-shot ICL. Bertsch

et al. (2025); Baek et al. (2025) report reduced sensitivity to demonstration selection, and Baek et al. (2025) highlight increased vulnerability to noisy examples in complex tasks. However, these works primarily focus on traditional classification-based ICL without exploring in-context CoT reasoning. Given the growing attention to the reasoning capabilities of LLMs and the demonstrated effectiveness with a large number of in-context demonstrations at test-time in enhancing model performance (Li et al., 2025), it becomes crucial to understand how many-shot CoT behaves under long-context settings. Our study provides a comprehensive evaluation of many-shot in-context CoT and investigates their unique characteristics, revealing deviations from the previously observed patterns in traditional many-shot ICL.

5.2 Chain-of-Thought

Prior studies have focused on modifying and enhancing the Chain-of-Thought (CoT) prompting paradigm to improve reasoning performance in large language models. Program of Thoughts (PoT) (Chen et al., 2023) introduces structured programming to represent the reasoning process more systematically. Tree-of-Thoughts (ToT) (Yao et al., 2023) proposes a tree-structured reasoning framework that enables the model to explore different reasoning paths. rStar-Math (Guan et al., 2025) decomposes complex reasoning problems and explores diverse reasoning trajectories using Monte Carlo Tree Search (MCTS) at test time, achieving significant improvements on mathematical reasoning benchmarks. In the meantime, only a few studies have explored the application of CoT in ICL. Dr.ICL (Luo et al., 2023) extends retrieval-augmented ICL to CoT prompting, demonstrating notable gains in mathematical reasoning tasks. However, these works typically operate under one-shot or few-shot ICL settings, leaving the potential of in-context CoT across extended context lengths largely underexplored. Our work addresses this gap by investigating how CoT prompting scales with increasing context length.

6 Conclusion

In this study, we provided a thorough investigation into the behavior of in-context CoT learning in extended context settings and its comparison with traditional ICL approaches for classification tasks. Our analysis uncovered several unique challenges

faced by in-context CoT learning, particularly its sensitivity to factors such as demonstration ordering and selection. This contrasts with prior findings in many-shot ICL, where such sensitivities are found to be less pronounced. Notably, we find that retrieving similar demonstrations does not enhance in-context CoT performance, diverging from established results in classification-based ICL.

Our empirical findings also demonstrate that task performance in many-shot in-context CoT is more influenced by the underlying data distribution more than task comprehension. Further experiments show that aligning in-context CoT demonstrations with LLMs' internal priors and learned reasoning trajectory can lead to more stabilized and consistent performance. Our work highlights the difference between in-context CoT with previous studies and the need for tailored strategies in leveraging in-context CoT learning, helping to lay the ground for further exploration of its potential and limitations.

Limitations

Due to the computational cost and performance limitations of LLMs in long in-context CoT reasoning, our study is limited to approximately 100 examples. While LLMs like Qwen 2.5 and LLaMA 3.1 can handle up to 131K and 128K context tokens, respectively, their performance in in-context CoT reasoning declines gradually beyond a certain threshold of context tokens, making exploring beyond 100 shots in this setting insignificant.

References

2021. [Benchmarking natural language understanding services for building conversational agents](#). In *Increasing Naturalness and Flexibility in Spoken Dialogue Interaction*, Lecture Notes in Electrical Engineering, pages 165–183. Springer.
- Rishabh Agarwal, Avi Singh, Lei M Zhang, Bernd Bohnet, Luis Rosias, Stephanie C.Y. Chan, Biao Zhang, Aleksandra Faust, and Hugo Larochelle. 2024. [Many-shot in-context learning](#). In *ICML 2024 Workshop on In-Context Learning*.
- Shengnan An, Zeqi Lin, Qiang Fu, Bei Chen, Nanning Zheng, Jian-Guang Lou, and Dongmei Zhang. 2023. [How do in-context examples affect compositional generalization?](#) In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11027–11052, Toronto, Canada. Association for Computational Linguistics.
- Jinheon Baek, Sun Jae Lee, Prakhar Gupta, Geunseob Oh, Siddharth Dalmia, and Prateek Kolhar. 2025. [Revisiting in-context learning with long context language models](#). *Preprint*, arXiv:2412.16926.
- Amanda Bertsch, Maor Ivgi, Emily Xiao, Uri Alon, Jonathan Berant, Matthew R. Gormley, and Graham Neubig. 2025. [In-context learning with long-context models: An in-depth exploration](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 12119–12149, Albuquerque, New Mexico. Association for Computational Linguistics.
- Iñigo Casanueva, Tadas Temčinas, Daniela Gerz, Matthew Henderson, and Ivan Vulić. 2020. [Efficient intent detection with dual sentence encoders](#). In *Proceedings of the 2nd Workshop on Natural Language Processing for Conversational AI*, pages 38–45, Online. Association for Computational Linguistics.
- Wenhu Chen, Xueguang Ma, Xinyi Wang, and William W. Cohen. 2023. [Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks](#). *Transactions on Machine Learning Research*.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reichiyo Nakano, Christopher Hesse, and John Schulman. 2021. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- Gilad Deutch, Nadav Magar, Tomer Natan, and Guy Dar. 2024. [In-context learning and gradient descent revisited](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 1017–1028, Mexico City, Mexico. Association for Computational Linguistics.
- Yiran Ding, Li Lyna Zhang, Chengruidong Zhang, Yuanyuan Xu, Ning Shang, Jiahang Xu, Fan Yang, and Mao Yang. 2024. [LongroPE: Extending LLM context window beyond 2 million tokens](#). In *Forty-first International Conference on Machine Learning*.
- Xinyu Guan, Li Lyna Zhang, Yifei Liu, Ning Shang, Youran Sun, Yi Zhu, Fan Yang, and Mao Yang. 2025. [rstar-math: Small llms can master math reasoning with self-evolved deep thinking](#). *Preprint*, arXiv:2501.04519.
- Chi Han, Qifan Wang, Hao Peng, Wenhan Xiong, Yu Chen, Heng Ji, and Sinong Wang. 2024. [LM-infinite: Zero-shot extreme length generalization for large language models](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3991–4008, Mexico City, Mexico. Association for Computational Linguistics.

688	Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. Measuring mathematical problem solving with the math dataset. <i>NeurIPS</i> .	743
689		744
690		
691		
692	Eduard Hovy, Laurie Gerber, Ulf Hermjakob, Chin-Yew Lin, and Deepak Ravichandran. 2001. Toward semantics-based answer pinpointing. In <i>Proceedings of the First International Conference on Human Language Technology Research</i> .	745
693		746
694		747
695		748
696		749
697	Yuxuan Huang, Yihang Chen, Haozheng Zhang, Kang Li, Huichi Zhou, Meng Fang, Linyi Yang, Xiaoguang Li, Lifeng Shang, Songcen Xu, Jianye Hao, Kun Shao, and Jun Wang. 2025. Deep research agents: A systematic examination and roadmap. <i>Preprint</i> , arXiv:2506.18096.	750
698		751
699		752
700		
701		
702		
703	Janak Kapuriya, Manit Kaushik, Debasis Ganguly, and Sumit Bhatia. 2025. Exploring the role of diversity in example selection for in-context learning. <i>Preprint</i> , arXiv:2505.01842.	753
704		754
705		755
706		756
707	Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. Large language models are zero-shot reasoners. In <i>Advances in Neural Information Processing Systems</i> .	757
708		
709		
710		
711	Tianle Li, Ge Zhang, Quy Duc Do, Xiang Yue, and Wenhui Chen. 2024. Long-context llms struggle with long in-context learning. <i>CoRR</i> , abs/2404.02060.	758
712		759
713		760
714	Yafu Li, Xuyang Hu, Xiaoye Qu, Linjie Li, and Yu Cheng. 2025. Test-time preference optimization: On-the-fly alignment via iterative textual feedback. <i>Preprint</i> , arXiv:2501.12895.	761
715		762
716		763
717		764
718	Bill Yuchen Lin, Abhilasha Ravichander, Ximing Lu, Nouha Dziri, Melanie Sclar, Khyathi Chandu, Chandra Bhagavatula, and Yejin Choi. 2024. The unlocking spell on base LLMs: Rethinking alignment via in-context learning. In <i>The Twelfth International Conference on Learning Representations</i> .	765
719		766
720		767
721		768
722		769
723		
724	Jiachang Liu, Dinghan Shen, Yizhe Zhang, Bill Dolan, Lawrence Carin, and Weizhu Chen. 2022. What makes good in-context examples for GPT-3? In <i>Proceedings of Deep Learning Inside Out (DeeLIO 2022): The 3rd Workshop on Knowledge Extraction and Integration for Deep Learning Architectures</i> , pages 100–114, Dublin, Ireland and Online. Association for Computational Linguistics.	770
725		771
726		772
727		773
728		774
729		775
730		776
731		777
732	Man Luo, Xin Xu, Zhuyun Dai, Panupong Pasupat, Mehran Kazemi, Chitta Baral, Vaiva Imbrasaitė, and Vincent Y Zhao. 2023. Dr.icl: Demonstration-retrieved in-context learning. <i>Preprint</i> , arXiv:2305.14128.	778
733		779
734		780
735		781
736		782
737	Costas Mavromatis, Balasubramaniam Srinivasan, Zhengyuan Shen, Jiani Zhang, Huzeifa Rangwala, Christos Faloutsos, and George Karypis. 2023. Which examples to annotate for in-context learning? towards effective and efficient selection. <i>Preprint</i> , arXiv:2310.20046.	783
738		
739		
740		
741		
742		
688	MetaAI. 2024. Introducing meta llama 3: The most capable openly available llm to date.	743
689		744
690		
691		
692	Sewon Min, Xinxin Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. 2022. Rethinking the role of demonstrations: What makes in-context learning work? In <i>Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing</i> , pages 11048–11064, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.	745
693		746
694		747
695		748
696		749
697		750
698		751
699		752
700		
701		
702		
703	Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shipppole. 2024. YaRN: Efficient context window extension of large language models. In <i>The Twelfth International Conference on Learning Representations</i> .	753
704		754
705		755
706		756
707		757
708		
709		
710		
711	Qwen, :, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jingren Zhou, and 25 others. 2025. Qwen2.5 technical report. <i>Preprint</i> , arXiv:2412.15115.	758
712		759
713		760
714		761
715		762
716		763
717		764
718	Charlie Victor Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2025. Scaling LLM test-time compute optimally can be more effective than scaling parameters for reasoning. In <i>The Thirteenth International Conference on Learning Representations</i> .	765
719		766
720		767
721		768
722		769
723		
724	Taylor Sorensen, Joshua Robinson, Christopher Rytting, Alexander Shaw, Kyle Rogers, Alexia Delorey, Mahmoud Khalil, Nancy Fulda, and David Wingate. 2022. An information-theoretic approach to prompt engineering without ground truth labels. In <i>Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 819–862, Dublin, Ireland. Association for Computational Linguistics.	770
725		771
726		772
727		773
728		774
729		775
730		776
731		777
732		778
733		
734		
735		
736		
737	Johannes von Oswald, Eyyvind Niklasson, Ettore Randazzo, João Sacramento, Alexander Mordvintsev, Andrey Zhmoginov, and Max Vladymyrov. 2023. Transformers learn in-context by gradient descent. <i>Preprint</i> , arXiv:2212.07677.	779
738		780
739		781
740		782
741		783
742		
724	Alex Wang, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R. Bowman. 2019. SuperGLUE: A stickier benchmark for general-purpose language understanding systems. In <i>Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada</i> , pages 3261–3275.	784
725		785
726		786
727		787
728		788
729		789
730		790
731		791
732		792
733		
734		
735		
736		
737	Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, brian ichter, Fei Xia, Ed H. Chi, Quoc V Le, and Denny Zhou. 2022. Chain of thought prompting elicits reasoning in large language models. In <i>Advances in Neural Information Processing Systems</i> .	793
738		794
739		795
740		796
741		797
742		

798 Zhiyong Wu, Yaoxiang Wang, Jiacheng Ye, and Ling-
799 peng Kong. 2023. Self-adaptive in-context learn-
800 ing: An information compression perspective for in-
801 context example selection and ordering. In *Proced-
802 ings of the 61st Annual Meeting of the Association for
803 Computational Linguistics (Volume 1: Long Papers)*,
804 pages 1423–1436, Toronto, Canada. Association for
805 Computational Linguistics.

806 Zhe Xu, Jiasheng Ye, Xiaoran Liu, Xiangyang Liu,
807 Tianxiang Sun, Zhigeng Liu, Qipeng Guo, Linlin
808 Li, Qun Liu, Xuanjing Huang, and Xipeng Qiu. 2025.
809 **Detectiveqa: Evaluating long-context reasoning on**
810 **detective novels.** *Preprint*, arXiv:2409.02465.

811 Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran,
812 Thomas L. Griffiths, Yuan Cao, and Karthik R
813 Narasimhan. 2023. Tree of thoughts: Deliberate
814 problem solving with large language models. In
815 *Thirty-seventh Conference on Neural Information
816 Processing Systems*.

817 Mo Yu, Tsz Ting Chung, Chulun Zhou, Tong Li, Rui
818 Lu, Jiangnan Li, Liyan Xu, Haoshu Lu, Ning Zhang,
819 Jing Li, and Jie Zhou. 2025. **Prelude: A benchmark**
820 **designed to require global comprehension and reason-**
821 **ing over long contexts.** *Preprint*, arXiv:2508.09848.

822 Zhuosheng Zhang, Aston Zhang, Mu Li, and Alex
823 Smola. 2023. **Automatic chain of thought prompting**
824 **in large language models.** In *The Eleventh Interna-*
825 *tional Conference on Learning Representations*.

826 A Correlating Factor Influencing 827 In-Context CoT Success

828 **Experiment Settings.** The experiments were
829 conducted on two datasets (i.e., BANKING77 and
830 GSM8K), both exhibit an overall trend of im-
831 provement as the number of in-context examples
832 increases, despite some fluctuations in GSM8K.
833 The evaluation is performed on the two smaller
834 LLMs (LLaMA 3.1, Qwen 2.5 (7B)). From the
835 test set, the first 300 instances are retrieved to
836 make a comparable evaluation to the testing ac-
837 curacy. In each retrieved instance, the latter
838 half of its question is masked and an external
839 LLM, LLaMA-3.3-70B-Instruct is used to gen-
840 erate continuations for the masked questions to
841 evaluate task understanding. The MCQA tasks are
842 constructed under two distinct conditions:

- 843 1. Task A: LLaMA 3.3 is instructed to avoid
844 the topic of the original task and generate
845 each continuations in entirely different do-
846 mains. This is to avoid LLM from locating
847 the ground-truth continuation easily with only
848 the option information. The clear differentia-
849 tion between the domains of options indicates
850 the task understanding. With more in-context
851 examples, the accuracy of Task A is expected
852 to continuously increase or plateau after a cer-
853 tain threshold. Examples of Task A are shown
854 in Appendix D.1.
- 855 2. Task B: LLaMA 3.3 is allowed to generate
856 continuations without the above restrictions,
857 leading to highly semantically similar option
858 sets with different expressions. Computing
859 the averaged pairwise cosine similarity among
860 the 4 options, Task B has a mean similarity
861 score of 0.726, which is significantly higher
862 than Task A of 0.471. Examples of Task B are
863 shown in Appendix D.2.

864 For human evaluation, a university graduate stu-
865 dent is invited to answer 120 sampled questions
866 provided with the in-context examples from the
867 training set. Humans can easily infer the task ob-
868 jectives and identify the correct option by aligning
869 it with the task domain of the training examples.
870 Results showed that humans achieved 93.33% ac-
871 curacy in Task A, significantly higher than their
872 performance of 46.67% accuracy in Task B.

873 **Result.** In the BANKING77 dataset, we observed
874 a clear positive trend in LLM performance as the

875 number of in-context examples increased. When
876 provided with in-context CoT with {question, CoT,
877 answer}, LLMs are prompted to predict the correct
878 continuation of masked questions. Surprisingly, as
879 shown in Table 1, Task B consistently exhibited
880 a higher spearman rank correlation with task ac-
881 curacy than Task A across all LLMs. This trend
882 contrasts with human performance.

883 To validate these findings in the reasoning task,
884 we conducted the same experiment on the GSM8K
885 dataset, creating analogous MCQA tasks under the
886 same conditions. Same conclusion is drawn with
887 Task B demonstrating a significantly stronger corre-
888 lation with task accuracy compared to Task A. All
889 the correlations with Task B are statistically signifi-
890 cant, with a p-value smaller than 0.05. It indicates
891 in-context CoT correlates more to data distribution
892 than task understanding.

		Task	Correlation	p-value
BANKING77	LLaMA	A	0.3376	0.4135
		B	0.8752	0.004
	Qwen	A	0.6412	0.0867
		B	0.9140	0.001
GSM8K	LLaMA	A	0.6190	0.1017
		B	0.7075	0.0496
	Qwen	A	0.6337	0.0916
		B	0.7933	0.0188

Table 1: Correlation between task accuracy and accu-
racy of the two constructed MCQA tasks.

893 B Prompt formatting and LLM 894 performance for each task

895 B.1 SuperGlue

896 We evaluate the Winograd Schema Challenge
897 (WSC) for coreference resolution, and the Choice
898 of Plausible Alternatives (COPA) for open-domain
899 commonsense causal reasoning. Both are format-
900 ted as a binary-label classification task. The prompt
901 for inference is presented in Figure 9 and 12, while
902 the evaluation result is shown in Figure 10 and 11
903 respectively.

904 B.2 TREC

905 We evaluate the Text REtrieval Conference (TREC)
906 Question Classification dataset with 50 fine class
907 labels. The prompt for inference is presented in
908 Figure 13, while the evaluation result is shown in
909 Figure 14.

Given a query, answer yes or no to the query.

The predicted answer must come from the demonstration examples with the exact format. The examples are as follows:

Question: In the sentence “{text₁ }”, does the pronoun ‘{span2_text₁}’ refer to {span1_text₁}?
Answer: {answer₁}

...

Question: In the sentence “{text_n }”, does the pronoun ‘{span2_text_n}’ refer to {span1_text_n}?
Answer: {answer_n}

Now predict the answer for the following query:

Question: In the sentence “{text_i }”, does the pronoun ‘{span2_text_i}’ refer to {span1_text_i}?

reply in the following format:

‘Answer: [yes | no]’

Figure 9: Prompt for WSC task




Figure 10: Performance on WSC

B.3 BANKING77

We evaluate the BANKING77 dataset with 77 fine-grained intents in the banking domain. The prompt for inference is presented in Figure 15, while the evaluation result is shown in Figure 16.

B.4 NLU

We evaluate the NLU dataset with 68 fine-grained intents in the conversational domain. The prompt for inference is presented in Figure 17, while the evaluation result is shown in Figure 18.

B.5 GSM8K

We evaluate the GSM8K dataset for grade school math word problems. The prompt for inference is presented in Figure 19, while the evaluation result is shown in Figure 20.

B.6 MATH

We evaluate the Mathematics Aptitude Test of Heuristics (MATH) dataset for mathematics competition problems, including the question types of counting_and_probability, prealgebra, geometry, precalculus, number_theory and algebra. The prompt for inference is presented in Figure 21, while the evaluation result is shown in Figure 22, 23, 24, 25, 26 and 27.

C Prompt for constructing MCQA Task A and B

C.1 Task A

The prompt to LLaMA-3.3-70B-Instruct for creating Task A is shown in Figure 28.

C.2 Task B

The prompt to LLaMA-3.3-70B-Instruct for creating Task A is shown in Figure 29.

910

911
912
913
914

915

916
917
918
919

920

921
922
923
924

925

926
927
928
929
930
931
932
933

934
935

936
937
938
939
940
941




Figure 11: Performance on COPA

942 **D Examples illustration of MCQA Task A
943 and B**

944 **D.1 Task A**

945 Two example illustrations in Task A, constructed
946 for BANKING77 and GSM8K, are shown in Figure
947 30. The option highlighted in bold is the correct
948 continuation of the incomplete question (Q).

949 **D.2 Task B**

950 Two example illustrations in Task B, constructed
951 for BANKING77 and GSM8K, are shown in Figure
952 31. The option highlighted in bold is the correct
953 continuation of the incomplete question (Q).

954 **E Prompt formatting for Task A and B**

955 The unified prompt for inference is presented in
956 Figure 32.

Answer in A or B.

The predicted answer must come from the demonstration examples with the exact format. The examples are as follows:

Premise: {premise₁}

Question: What is the {question₁} for this?

Options:

A. {choice1₁}

B. {choice2₁}

Answer: {answer₁}

...

Premise: {premise_n}

Question: What is the {question_n} for this?

Options:

A. {choice1_n}

B. {choice2_n}

Answer: {answer_n}

Now predict the answer for the following query:

Premise: {premise_i}

Question: What is the {question_i} for this?

Options:

A. {choice1_i}

B. {choice2_i}

reply in the following format:

'Answer: [A | B]'

Figure 12: Prompt for COPA task

Given a question, predict the label of the question. You can only make predictions from the following categories: {LIST_OF_CATEGORIES}

Please predict the label of the FINAL question with the provided demonstration example queries as follows:

question: {question₁}

label: {label₁}

...

question: {question_n}

label: {label_n}

Now predict the answer for the following query:

question: {question_i}

reply in the following format:

'label: [category_name]'

Figure 13: Prompt for TREC task




Figure 14: Performance on TREC

Given a question, predict the label of the question. You can only make predictions from the following categories:
{LIST_OF_CATEGORIES}
 Please predict the intent category of the FINAL query with the provided demonstration example queries as follows:

```

service query: {question1}
intent category: {label1}
...
service query: {questionn}
intent category: {labeln}
```

Now predict the intent category for the following query:

```
service query: {questioni}
```

reply in the following format:
 'intent category: [category_name]'

Figure 15: Prompt for BANKING77 task




Figure 16: Performance on BANKING77

Given a question, predict the label of the question. You can only make predictions from the following categories:
{LIST_OF_CATEGORIES}

Please predict the intent category of the FINAL utterance with the provided demonstration example queries as follows:

utterance: {question₁}
intent category: {label₁}

...

utterance: {question_n}
intent category: {label_n}

Now predict the intent category for the following utterance:

utterance: {question_i}

reply in the following format:

'intent category: [category_name]'

Figure 17: Prompt for NLU task




Figure 18: Performance on NLU

In the end of the response, add a summary 'The answer is [answer].'

Q: {question₁}
A: {CoT₁} {answer₁}

...

Q: {question_n}
A: {CoT_n} {answer_n}

Q: {question_t}
A: Let's think step by step.

Figure 19: Prompt for GSM8K task




Figure 20: Performance on GSM8K

Write a response that appropriately completes the request and wrap the final answer inside `\boxed{}`.

```

Problem: {question_1}
Solution: {CoT_with_answer_1}
...
Problem: {question_n}
Solution: {CoT_with_answer_n}

### Problem: {question_t}
### Solution: Let's think step by step.

```

Figure 21: Unified prompt for MATH task




Figure 22: Performance on counting_and_probability




Figure 23: Performance on prealgebra




Figure 24: Performance on geometry




Figure 25: Performance on precalculus




Figure 26: Performance on number_theory




Figure 27: Performance on algebra

BANKING77:

Task: Task: Given an incomplete customer query, generate 6 unique and diverse continuations to complete the question.
Ensure:

- The continuations avoid topics related to banking or customer service queries to ensure diversity.
- Each continuation explores different contexts or domains to reflect a variety of possibilities. Please diversify your generation with the provided example.

Input:

Q: {masked_question}

Output:

1. {groundtruth}
- 2.
- 3.
- 4.
- 5.

GSM8K:

Task: Given an incomplete question, generate 6 unique and diverse continuations to complete the question. Ensure:

- The continuations avoid topics related to math or grade school level difficulty to ensure diversity.
- Each continuation explores different contexts or domains to reflect a variety of possibilities. Please diversify your generation with the provided example.

Input:

Q: {masked_question}

Output:

1. {groundtruth}
- 2.
- 3.
- 4.
- 5.

Figure 28: Prompt for constructing Task A

Given an incomplete math question, generate a 6 version of continuation to complete the question. Ensure that the number of tokens in your completion is approximately equal to the number of tokens in the provided incomplete question.

Input:

Q: {masked_question}

Output:

1. {groundtruth}
- 2.
- 3.
- 4.
- 5.

Figure 29: Unified prompt for constructing Task B

BANKING77:

Q: When traveling, can I auto

Options:

- A. track my daily expenses and stay within a set budget using a travel app?
- B. navigate through unfamiliar cities using augmented reality maps?
- C. unlock my hotel room door using a digital key on my smartwatch?
- D. top-up my card at certain times?**

GSM8K:

Q: Every day, Wendi feeds each of her chickens three cups of mixed chicken feed, containing seeds, mealworms and vegetables to help keep them healthy. She gives the chickens their feed in three separate meals. In the morning, she gives her flock of chickens 15

Options:

- A. automated feeding system that dispenses the exact amount of feed at each meal, reducing waste and saving time. What are the potential advantages and disadvantages of using an automated feeding system for Wendi's flock?
- B. varieties of vegetables, such as kale and carrots, to supplement their diet and provide essential nutrients. How can Wendi incorporate these vegetables into the chickens' feed in a way that is both cost-effective and efficient?
- C. cups of feed. In the afternoon, she gives her chickens another 25 cups of feed. How many cups of feed does she need to give her chickens in the final meal of the day if the size of Wendi's flock is 20 chickens?**
- D. local farm-to-table initiative, where she sells the eggs produced by her chickens to nearby restaurants and cafes. What are some key marketing strategies that Wendi could use to promote her farm-to-table business and attract new customers?

Figure 30: Examples illustration of Task A. The correct answer is highlighted in bold.

BANKING77:

Q: If I request that my funds

Options:

- A. be withdrawn, what are the minimum requirements?
- B. be held, what currencies do you use?**
- C. be converted, what is the exchange rate?
- D. be deposited, what is the maximum limit?

GSM8K:

Q: Toulouse has twice as many sheep as Charleston. Charleston has 4 times as many sheep as

Options:

- A. Seattle. What is the combined total of sheep in Toulouse, Charleston, and Seattle if Seattle's sheep population is 10?
- B. Seattle. How many sheep do Toulouse, Charleston, and Seattle have together if Seattle has 20 sheep?**
- C. Seattle. How many sheep are there in total in Seattle, Charleston, and Toulouse if Seattle has 15 sheep?
- D. Seattle. If Seattle is home to 5 sheep, what is the total number of sheep in Charleston, Toulouse, and Seattle altogether?

Figure 31: Examples illustration of Task B. The correct answer is highlighted in bold.

BANKING77:

Your task is to choose the best option from the four provided that completes the question.

Do NOT solve or answer the question; ONLY respond with the correct option label (A, B, C, or D).

```
service query: {question1}  
intent category: {label1}  
...  
service query: {questionn}  
intent category: {labeln}
```

```
Final query: {mcqa_question}  
A. {optionA}  
B. {optionB}  
C. {optionC}  
D. {optionD}
```

GSM8K:

Your task is to choose the best option from the four provided that completes the question.

Do NOT solve or answer the question; ONLY respond with the correct option label (A, B, C, or D).

```
Question: {question1}  
Answer: {CoT1} {answer1}  
...  
Question: {questionn}  
Answer: {CoTn} {answern}
```

```
Final Question: {mcqa_question}  
A. {optionA}  
B. {optionB}  
C. {optionC}  
D. {optionD}
```

Figure 32: Unified prompt for Task A and B